

3**. Serverless Data Processing Pipeline**

Objective: Build a serverless pipeline for processing data (e.g., log processing or ETL jobs).

Approach:

- **Data Ingestion:** Use AWS services like S3 or Kinesis to ingest data.
- **Processing:** Create Lambda functions to process the ingested data.
- **Storage:** Store the processed data in an appropriate AWS service, like S3 or DynamoDB.
- **Monitoring:** Set up CloudWatch to monitor the pipeline's performance and to log any issues.

Goal: Learn to build a serverless data processing pipeline, understanding the flow of data through various AWS services.

1. Create S3 bucket:

The screenshot shows the 'Create bucket' page in the AWS Management Console. The breadcrumb navigation at the top reads 'Amazon S3 > Buckets > Create bucket'. The main heading is 'Create bucket' with an 'Info' link. Below this, a sub-header states 'Buckets are containers for data stored in S3.' The 'General configuration' section is active. It features a dropdown for 'AWS Region' set to 'US East (N. Virginia) us-east-1'. Under 'Bucket type', the 'General purpose' option is selected with a radio button, while 'Directory - New' is unselected. The 'General purpose' description notes it is recommended for most use cases and allows a mix of storage classes. The 'Bucket name' field contains 'pipelinebuckrt'. A note below the field states the name must be unique and follow naming rules, with a link to 'See rules for bucket naming'. At the bottom, there is a section for 'Copy settings from existing bucket - optional' with a 'Choose bucket' button and a format example 's3://bucket/prefix'.

Amazon S3 > Buckets > Create bucket

Create bucket [Info](#)

Buckets are containers for data stored in S3.

General configuration

AWS Region

US East (N. Virginia) us-east-1 ▼

Bucket type [Info](#)

☒ **General purpose**

Recommended for most use cases and access patterns. General purpose buckets are the original S3 bucket type. They allow a mix of storage classes that redundantly store objects across multiple Availability Zones.

☐ **Directory - New**

Recommended for low-latency use cases. These buckets use only the S3 Express One Zone storage class, which provides faster processing of data within a single Availability Zone.

Bucket name [Info](#)

pipelinebuckrt

Bucket name must be unique within the global namespace and follow the bucket naming rules. [See rules for bucket naming](#) [↗](#)

Copy settings from existing bucket - optional

Only the bucket settings in the following configuration are copied.

Choose bucket

Format: s3://bucket/prefix

2. Create Lambda Function:

[Lambda](#) > [Functions](#) > [Create function](#)

Create function Info

Choose one of the following options to create your function.

☒ **Author from scratch**
Start with a simple Hello World example.

☐ **Use a blueprint**
Build a Lambda application from sample code and configuration presets for common use cases.

☐ **Container image**
Select a container image to deploy for your function.

Basic information

Function name
Enter a name that describes the purpose of your function.

Use only letters, numbers, hyphens, or underscores with no spaces.

Runtime Info
Choose the language to use to write your function. Note that the console code editor supports only Node.js, Python, and Ruby.

Python 3.12 ▼ ↻

Architecture Info
Choose the instruction set architecture you want for your function code.

☒ x86_64
☐ arm64

Permissions Info
By default, Lambda will create an execution role with permissions to upload logs to Amazon CloudWatch Logs. You can customize this default role later when adding triggers.

▼ Change default execution role

Execution role
Choose a role that defines the permissions of your function. To create a custom role, go to the [IAM console](#).

☐ Create a new role with basic Lambda permissions
☒ Use an existing role

3. Code was placed to transform the received data to upper case and it triggers when the file is uploaded to s3 bucket

Successfully created the function processinglambda. You can now change its code and configuration. To invoke your function with a test event, choose "Test".

Function URL Info

[Code](#) [Test](#) [Monitor](#) [Configuration](#) [Aliases](#) [Versions](#)

Code source Info

[Upload from](#)

File Edit Find View Go Tools Window **Test** Deploy Changes not deployed

Go to Anything (Ctrl-F)

Environment


```
1 import boto3
2 import urllib.parse
3 import os
4
5 def lambda_handler(event, context):
6     s3_client = boto3.client('s3')
7
8     # Get bucket name and object key from the S3 event
9     bucket_name = event['Records'][0]['s3']['bucket']['name']
10    object_key = urllib.parse.unquote_plus(event['Records'][0]['s3']['object']['key'], encoding='utf-8')
11
12    # Extract the filename from the object key
13    filename = os.path.basename(object_key)
14
15    # Define destination key (output folder and file name)
16    destination_key = 'output/' + filename
17
18    # Get the file from S3
19    file_obj = s3_client.get_object(Bucket=bucket_name, Key=object_key)
20    file_content = file_obj['Body'].read().decode('utf-8')
21    s3_client.put_object(Bucket=bucket_name, Key=destination_key, Body=file_content)
22
23    # Convert content to uppercase
24    upper_content = file_content.upper()
25
26    # Upload the modified content back to S3
27    s3_client.put_object(Bucket=bucket_name, Key=destination_key, Body=upper_content)
28
29
30
31    return {
32        'statusCode': 200,
33        'body': 'File processed and uploaded successfully'
34    }
```

4. Add trigger

[Lambda](#) > **Add trigger**

Add trigger

Trigger configuration [Info](#)

 **S3**
aws asynchronous storage

Bucket
Choose or enter the ARN of an S3 bucket that serves as the event source. The bucket must be in the same region as the function.

Bucket region: us-east-1

Event types
Select the events that you want to have trigger the Lambda function. You can optionally set up a prefix or suffix for an event. However, for each bucket, individual events cannot have multiple configurations with overlapping prefixes or suffixes that could match the same object key.

All object create events

Prefix - optional
Enter a single optional prefix to limit the notifications to objects with keys that start with matching characters.

Suffix - optional
Enter a single optional suffix to limit the notifications to objects with keys that end with matching characters.

5.Upload file to s3 bucket

[Amazon S3](#) > [Buckets](#) > [pipelinebuckrt](#) > **Upload**

Upload Info

Add the files and folders you want to upload to S3. To upload a file larger than 160GB, use the AWS CLI, AWS SDK or Amazon S3 REST API. [Learn more](#)

Drag and drop files and folders you want to upload here, or choose **Add files** or **Add folder**.

Files and folders (1 Total, 32.0 KB)
All files and folders in this table will be uploaded.

☐

Name

Folder

<input type="checkbox"/>	random.txt	-
--------------------------	------------	---

Destination Info

Destination

s3://pipelinebuckrt

► Destination details

5. Here is the log report from the cloud watch that lambda function was triggered successfully

CloudWatch × [CloudWatch](#) > [Log groups](#) > [/aws/lambda/1stserverless](#) > **All events**

Log events
You can use the filter bar below to search for and match terms, phrases, or values in your log events. [Learn more about filter patterns](#)

Clear

1m

30m

1h

12h

Custom

Local timezone

Display

Timestamp	Message	Log stream name
2024-02-25T11:01:53.777+05:45	INIT_START Runtime Version: python:3.12.v19 Runtime Version ARN: arn:aws:lambda:us-east-1::runtime:a79ae1d0439e8e7a1dc8b45a221e8f0b0bc495c3ff8-	2024/02/25/[SLATEST]ja467c802ac44b7aabbdb6d416c17da
2024-02-25T11:01:53.860+05:45	START RequestId: 2d2abdb-d2cf-4738-8ff6-affeec35287 Version: \$LATEST	2024/02/25/[SLATEST]ja467c802ac44b7aabbdb6d416c17da
2024-02-25T11:01:53.860+05:45	END RequestId: 2d2abdb-d2cf-4738-8ff6-affeec35287	2024/02/25/[SLATEST]ja467c802ac44b7aabbdb6d416c17da
2024-02-25T11:01:53.860+05:45	REPORT RequestId: 2d2abdb-d2cf-4738-8ff6-affeec35287 Duration: 7.25 ms Billed Duration: 8 ms Memory Size: 128 MB Max Memory Used: 34 MB Init Du.	2024/02/25/[SLATEST]ja467c802ac44b7aabbdb6d416c17da
2024-02-25T11:02:03.930+05:45	START RequestId: c52620be-153b-4ed1-a56d-84ef3343a308 Version: \$LATEST	2024/02/25/[SLATEST]ja467c802ac44b7aabbdb6d416c17da
2024-02-25T11:02:03.931+05:45	END RequestId: c52620be-153b-4ed1-a56d-84ef3343a308	2024/02/25/[SLATEST]ja467c802ac44b7aabbdb6d416c17da
2024-02-25T11:02:03.931+05:45	REPORT RequestId: c52620be-153b-4ed1-a56d-84ef3343a308 Duration: 1.66 ms Billed Duration: 2 ms Memory Size: 128 MB Max Memory Used: 34 MB	2024/02/25/[SLATEST]ja467c802ac44b7aabbdb6d416c17da
2024-02-25T11:02:05.164+05:45	START RequestId: b0c977e7-f5d9-4f46-8d0c-e0aabcfc62c12 Version: \$LATEST	2024/02/25/[SLATEST]ja467c802ac44b7aabbdb6d416c17da
2024-02-25T11:02:05.160+05:45	END RequestId: b0c977e7-f5d9-4f46-8d0c-e0aabcfc62c12	2024/02/25/[SLATEST]ja467c802ac44b7aabbdb6d416c17da
2024-02-25T11:02:05.160+05:45	REPORT RequestId: b0c977e7-f5d9-4f46-8d0c-e0aabcfc62c12 Duration: 1.05 ms Billed Duration: 2 ms Memory Size: 128 MB Max Memory Used: 34 MB	2024/02/25/[SLATEST]ja467c802ac44b7aabbdb6d416c17da
2024-02-25T11:35:47.585+05:45	INIT_START Runtime Version: python:3.12.v19 Runtime Version ARN: arn:aws:lambda:us-east-1::runtime:a79ae1d0439e8e7a1dc8b45a221e8f0b0bc495c3ff8-	2024/02/25/[SLATEST]ja467351fc5854d1c9736d279defc779
2024-02-25T11:35:47.585+05:45	START RequestId: 64706f55-0c31-4d2d-8a38-b0b0f96a3bb1 Version: \$LATEST	2024/02/25/[SLATEST]ja467351fc5854d1c9736d279defc779
2024-02-25T11:35:47.670+05:45	END RequestId: 64706f55-0c31-4d2d-8a38-b0b0f96a3bb1	2024/02/25/[SLATEST]ja467351fc5854d1c9736d279defc779
2024-02-25T11:35:47.670+05:45	REPORT RequestId: 64706f55-0c31-4d2d-8a38-b0b0f96a3bb1 Duration: 2.44 ms Billed Duration: 3 ms Memory Size: 128 MB Max Memory Used: 34 MB Init Du.	2024/02/25/[SLATEST]ja467351fc5854d1c9736d279defc779

6. Here is the query that our lower case data was successfully transformed to upper case

SQL query

Amazon S3 Select supports only the SELECT SQL command. Using the S3 console, you can extract up to 40 MB of records from an object that is up to 128 MB in size. To work with larger files or more records, use the AWS CLI, AWS SDK, or Amazon S3 REST API. For more complex SQL queries, use [Amazon Athena](#).

1

/* To create reference point for writing SQL queries, you can display the first 5 records of input data by running the following SQL query: SELECT * FROM s3object s LIMIT 5 */

2

SELECT * FROM s3object s LIMIT 5

SQLLn 1, Col 1Errors: 0Warnings: 0

Query results

Query results are not available after you choose Close or navigate away. Choose [Download results](#) to download a copy of the following query results.

Status

Successfully returned 1 record in 2948 ms

Bytes returned: 32769 B

1

```
EOGPHHFECEQROPOCVCZKSADE13CQFQVBDCTBOADTARZHQ3QOHTSRQRYWEZKULAW3LHWAOK1BA8YHPLUYFADGKYGABRTREBQROCSWUCHYAGOL3PZOVOKTVYDKGXVQ1Y1E2SFKLMPHIEFFDQZRES3TOXP3KQUNRYT3AOVAORAPHHATRES1YNSPSPFUNNVOPUD3TNHYNIEPCC320CKQWDEKXAUYNMKL
3OYXMKLYNATHYBOKD2ZFALRAPHDSVALHQWYEPNGS3ARJ2NT3HEKXB7ASRVINTNBS32QORLZEPGGFCBAZ3HBYHTWPC3LTHGADVDH3VAL3CTKRNCHTYLHMSVAGULYLYHBP3LINTHOLD3LJVLKXNTWPHIDEP3JUNCS3VPMXYSKQZGKZ3MTOGKQBS33TOBPF3CLJCKOXYLHPPYTHCLNFM3Q
TNV7ZB333PCOP3DGLNTEGAL3KDEK3KLCNS4BYX3ZNV3G3BUNQ3E333DGL3PMAYRTS3ZTH3OYQ3HAG3C3K3ONB33ITUD3WV3QL3RAD3Q3K3PT3OT3Z3NET3Q3HT3ONCE3Z3B3N3E3Y3M3K3NL3T3ON333Q3P3Q3CH3E3G3FT3O3K333AT3Y3P3OT3O3NT3K3U3Z3G3H3AL3K3E3L3W3P3C3TT3H3Q3FF3O3R3CE3E3VT3DL3V
GHUEVZG3E3L3T3P3O3Q3Z3B3M3C3UD3Y3L3H3NT33CV3B3G3M3L3O3CP3H3T3H3F3O3K3R3C3Y3V3N3AL3Q3B3SP3Y3H3N3D3C3Y3O3M3C3E3G3V3K3P3R3E3D3Q333V3B3T33Y33T3F3Q3D3K3O3D3H3T3R3H3C3G3S3F3D3G3B3O3R3F3Q3P3K3Z3N3AR3AD3E3T3R3U3Y3K3C3P3V3O3K3P3A3T3I3Z3E3B3P3L3Z3N3S3Y3E3R3O3Z33Q3H3T3K3R3G3K3L3Q3H3C3Y3O3B3T3R3A3W3H3P3T3K3U3Z3N3T
H3B3F3N3CTN3D3K3E3Y3A3P37233L3E3N3D3A3B3T3Y3P3F3D3S3A3H3T3C3E33F3A3D3Y3Q3Z3B3M3C3UD3Y3L3H3NT33G3O3P3G33Z33B3A3P3Y3T3H3P3C3Z3B3U3S3P3D3V33W3N3T3G3A3T3I3Z3L3F3Q3D3CT3U3G3P3G3H3L3D3A3G3F3A3K3E3S3D3G3U3P3B3L3S3C3S3E3Q3T3W3A3B3H3C3K3G3P3T3D3Z3A3Y3A3F3T3U3Q3E3H3P3U3A3H3B3Z33H3T3H3T3E3C3O3T3S3F3P3S3C3P
Q3Q3H3B3U3B3Z3C3N3L3E3Y3O3L3H333T3U3A3F3C3Q3P3F3U3N3E3F3B3H3P3N3Y3E3M3T3Z3I3E3D3N3Q3C3A3B3V3E3D3R3G3H3Y3Y3E3W3H3Q3O3Z3E3F3E3N3Q3L3Q3U3F3A3H3K3Z3H3B333P3C3H3A3L3H3K3O3E3P3H3K3B3Z3O3V3T3H3K3Z3H3Y3E3U3F3Y3L3Y3A3R3O3O3V3K3Q3N3A3T3E3K3P3F3L3L3N3N3E3H3X3H3C3N3E3F3S3C3Q33T3H3N3Y3T3E3D3A3N3Y3P3A3P3Z3H3G3T3Z3H3Y3N3B3
X3A3A3H3D3P3V3B3G3A3R3F3353K3Z3S3B3E3H3C3Z3C3M3E3B3Y3T3S3D3P3C3Q3D3E3J3D3Y3W3H3L3K3D3N3R3G3T3Y3V3U3V3U3B3L3B3O3L3D3C3K3Z3G3H3K3J3U3S3A3E3H3N3G3T3P3G3W3O3T3S3H3T33T3O3L3T3G3L3V3H3R3L3P3G3E3Q3L3T3K3X3O3A3T3S3U3P3V3Z3G3P3Z3N3H3Q3Q3O3R3F3E3Q3Y3T3Z3P3W3V3L3A3N3B3T3T3A3C3N3F3A3E3T3P3E3Z33F3R3E3N3G3R3Q3K3D3L3S3Y3Z
V3E3Z3N3P3Y3W3K3L3L3L3U3G3T3O3T3E3T3E3T3A3T3C3P3H3E3G3T3Y3E3N3B3U3G3A3B3L3B3H3K3Z3N3U3F3Q3Q3Y3C3H3B3W3O3T3Q3T3B3K3A3T3Y3T3E3A3B3F3S3T3B3U3K3A3F3I3O3Z3L3H3A3Y33E3N3E3G3H3D3U3T3A3C3Y3T3E3R3H3A3L3A3T3O3C3P3H3U3V3W3K3Y3O3N3G3V3E3Y3H3Z3K3E3Y3H3L3A3P3S3E3D3G3V3D3C3U3C3T3O3T3N3A3G3S3H3C3T3P3L3S3L3W3A3H3W3Y3N3S3D3R3T3Q3N
H3A3F3V3E3L3M3E3L3I3U3H3P3U3R3D3G3H3P3F3G3U3R3Z3L3P3G3H3N3D3C3L3P3Z3Z3D3W3Q3V3K3N3D3C3P3A3K3G3H3K3P3Q3Y3E3P3H3D3C3Q3K3L3A3D3Q3A3N3K3U3A3V3A3H3P3C3H3T3A3L3G3A3M3A3N3K3U3O3T3E3Q3V3Z3K3H3F3K3U3E3Y3L3W3A3G3Q3R3N3Q3L3I3P3B3B3Y3B333B3E3U3T3R3E3K3O3O3R3P3U3H3P3K3L3Y3L3A3L3M3Q3N3T3H3Y3D3C3K3E3Z3E3N3K3S3D3M3P3Y3E3K3C3B3C3L
L3O3H3F3A3Z3E3E3N3H3E3V3L3U3Z3A3P3C3A3H3W3H3Y3D3L3C3V3D3S3P3G3Z3E3D3W3Y3L3D3S3F3D3H3K3C3K3Q3Y3U3D3H3N3B3T3E3S3B3H3E3Q3E3C3K3H3A3C3B3P3W3Q3V3O3L3D3U3T3H3Q3V3P3O3K3C3A3P3C3F3R3K3E3T3F3W3H3G3L3K3A3D3C3Y3F3L3H3C3E3F3E3S3H3K3E3C3H3A3P3H3E3F33H3T3R3D3F33H3Y3L3A3A3P3A3Z3W3E3F3P3Z3H3Z3D3X3V3A3P
N3H3P3Z3G3Y3U3Z3D3E3X3G3A3V3E3G3N3Y3K3Y3V3A3C3D3S3Z3I3G3Z3U3B3K3Z3Q3L3L3O3T3P3Z3U3F3X3H3H3R3B3V3Y3Q3A3L3H3L3H3P3H3R3G3E3K3B3X3Z3G3T3Y3E3D3E3D3K3T3Y3E3X3Z3Q3O3U3D3Y3V3A3Z3E3T3P3Z3U3S3Q3B3T3O3G3R3V3Z3U3Z3U3M3U3H3Q3B3L3H3C3P3H3G3V3A3H3A3K3T3U3E3L3Y3P3Z3V3D3C3E3R3G3P3Z3K3I3O3N3K3B3C3A3F33W3P3A3D3C3Y3U3B3Q3Z3K3
Q3T3S3T3Y3Q3D3C3Y3L3K3C3L3Q3O3E3D3A3R3D3T3H3K3N3K3W3A3U3S3C3H3K3N3C3B3K3T3F3G3I3A3N3G3D3Y3P3H3A3E3C3G3G3O3X3T3V3H3S3U3F3D3C3U3P3Q3C3O3P3O3K3W3A3T3F3U3G3O3N3Q3K3B3R3N3D3F32E3G3T3C3Q3Z3O3F3T3O3Y3A3B3H3T3E3P3A3Y3L3L3O3U3S3D3P3F3C3H3A3P3G3H3K3C3X3A3A3R3A3B3H3H3S3U3X3Z3W3K3E3Y3E3G3Q3B3U3T3O3C3H3P3G3O3N3Q3D3B3T
B3O3P3W3E3D3E3F3T3E3P3O3E3F3A3T3Y3P3B3Z3T3Y3K3W3A3T3O3P3G3H3E3L3W3S3C3T3F3U3L3P3K3H3G3T3A3A3N3K3D3O3N3A3C3H3A3G3L3A3U3R3D3L3Q3H3P3C3Q3H3S3U3Y3G3E3Z3B3D3H3T3E3Z3O3N3Q3H3Z3A3L3T3B3A3Z3P3A3T3H3D3O3B3F3A3P3B3C33H3Y3A3H3K3T3U3P3K3L3D3C3F3P3H3A3K3E3Z3O3T3Y3P3B3C3D3W3C3D3P3Y3A3C3H3H3Y3Q3Z3L3G3U
T3S3Z3Q3A3E3X3P3C3Q3P3Y3A3C3V3F3H3T3X3Y3I3Q3Y3K3T3Y3H3Z3B3S3D3C3P3E3C3A3H3Q3T3B3Z3K3B3E3Y3F3G3I3Z3T3H3D3B3H3T3X3P3O3T3A3D3S3G3K3L3S3Y3P3V3K3R3O3F3G3H3L3P3A3Y3A3Q3V3O3M3B3L3Z3U3E3Q3U3K3P3I3F3Z3Q3Y3E3W3H3F3A3T3F3S3K3I3N3S3Y3O3N3H3Q3U3R3H3P3H3K3S3U3R3L3G3A3V3D3C3R3U3T3D3T3O3U3R3E3Z3Z3H3Y3Z3I3Z3H3T3H3T3Y3E3C3P3T3
M3Q3H3T3Y3P3M3T3T3Q3H3C3L3A3H3D3V3D3K3Q3S3V3B3N3C3H3E3S3U3A3V3I3B3Z3U3P3H3D3Q3G3F3T3A3N3T3O3R3C3Q3T3E3Z3Q3T3Q3L3T3A3Q3L3T3D3N3P3Y3G3F3S3Z3C3L3A3B3E3Y3T3O3L3A3P3A3L3G3U3G3H3A3G3G3P3Y3T3H3A3Z3D3Z3A3G33Y3L3H3D3H3Z3D3C3Y3H3V3K3A3Q3H3S3T3K3C3I3E3X3T3B3I3T3E3P3M3C3S3Y3B3D3T3L3E3T3A3P3A3B3C33A3E3H3S3G3Q3S3Q3D3W3H3E3Y3L
```